

Why and How to Roll Your Own `std::function` Implementation

Tom Poole - CppCon 2018

This talk

The supporting material is available on GitHub:

<https://github.com/tpoole/presentations>

I'm a maintainer of the open source, cross-platform, application development framework JUCE.

This talk is motivated by the addition of a `std::function` replacement to JUCE.

- Cross-platform compatibility
- JUCE is focussed on audio applications

But first, a `std::function` recap...

What is std::function?

A type-erasing wrapper containing a callable thing with a specific signature

```
int addOneFunction (int x) {  
    return x + 1;  
}  
  
struct AddOneStruct {  
    int operator() (int x) const {  
        return x + 1;  
    }  
};  
  
std::function<int(int)> f;  
f = addOneFunction;  
f = AddOneStruct();  
f = [](int x) { return x + 1; };
```

Callable things:

- Function pointers
- Objects with a call operator
- Lambda functions

Why is that useful?

`std::function` has transformed the way many programs are written

- Asynchronous callbacks

`std::function` is a good fit for audio processing operations

- A lot of audio effects are successive transformations of arrays of audio samples
- Each stage has the same signature (input samples -> output samples)
- Swappable arrays of callable objects a good way to change behaviour dynamically

Example: An oscillator class which takes a `std::function` from a user to determine the periodic part of a waveform

Cross platform compatibility

JUCE is cross platform: desktop, mobile, embedded

A decent number of systems lack a standard library containing `std::function`

- Toolchains for embedded systems
- OS X < 10.7

JUCE users want to be able to use the same code everywhere

Realtime audio safety

Processing live audio is **merciless**.

The process is typically:

- You supply the audio device with a callback, with a pointer to some memory and a size
- The callback is called from the audio thread
- You fill the memory with some audio samples
- A short time later the memory is read by the sound card

If you don't finish writing to the memory you will hear a glitch.

Predictable runtime

Glitches are **bad!**

- People will not use your audio software
- Potential to cause permanent hearing damage

Golden rule of audio programming:

- **Do not do anything on the audio thread which can take an indeterminate amount of time**
 - > No locks
 - > No memory allocations

When does `std::function` allocate?

`std::function` can avoid allocations by using a small bit of stack space in each object

- Small buffer optimisation (the same principle as `std::string`)

If a callable thing is bigger than the stack space then `std::function` will get some memory from the heap

- The size of the stack space is implementation defined

Other std::function alternatives

Game developers are also worried memory allocations

SG14 (GameDev & low latency ISO C++ working group):

- `inplace_function`

Facebook's Folly library (see last year's CppCon talk by Sven Over):

- `folly::Function`

A move-only alternative to std::function which also addresses const-correctness

A basic implementation

Define a template class, then pass the function signature you want into that template parameter

```
template <typename>  
class function;  
  
template <typename Result, typename... Arguments>  
class function<Result (Arguments...)>
```

Type-erasing

```
using invokePtr_t = Result(*) (void*, Arguments&&...);  
using createPtr_t = void(*) (void*, void*);  
using destroyPtr_t = void(*) (void*);  
  
invokePtr_t invokePtr;  
createPtr_t createPtr;  
destroyPtr_t destroyPtr;  
  
std::unique_ptr<char[]> storage;
```

```
template <typename Functor>  
static Result invoke (Functor* f, Arguments&&... args)  
{  
    return (*f)(std::forward<Arguments>(args)...);  
}  
  
template <typename Functor>  
static void create (Functor* destination, Functor* source)  
{  
    new (destination) Functor (*source);  
}  
  
template <typename Functor>  
static void destroy (Functor* f)  
{  
    f->~Functor();  
}
```

Constructor, destructor, call operator

```
template <typename Functor>
function (Functor f)
    : invokePtr (reinterpret_cast<invokePtr_t> (invoke<Functor>)),
      createPtr (reinterpret_cast<createPtr_t> (create<Functor>)),
      destroyPtr (reinterpret_cast<destroyPtr_t> (destroy<Functor>)),
      storage (new char[sizeof (Functor)])
{
    createPtr (storage.get(), std::addressof (f));
}

~function()
{
    destroyPtr (storage.get());
}

Result operator() (Arguments&&... args) const
{
    return invokePtr (storage.get(), std::forward<Arguments> (args)...);
}
```

Copying, assigning and null/empty state

```
function() = default;

function (const function& other)
{
    if (other.storage != nullptr)
    {
        invokePtr  = other.invokePtr;
        createPtr  = other.createPtr;
        destroyPtr = other.destroyPtr;

        storageSize = other.storageSize;
        storage.reset (new char[storageSize]);

        createPtr (storage.get(), other.storage.get());
    }
}
```

```
size_t storageSize;
std::unique_ptr<char[]> storage;
```

```
function& operator= (function const& other)
{
    if (storage != nullptr)
    {
        destroyPtr (storage.get());
        storage.reset();
    }

    if (other.storage != nullptr)
    {
        invokePtr = other.invokePtr;
        createPtr = other.createPtr;
        destroyPtr = other.destroyPtr;

        storageSize = other.storageSize;
        storage.reset (new char[storageSize]);

        createPtr (storage.get(), other.storage.get());
    }

    return *this;
}
```

Benchmark against std::function

```
int addOne (int x)
{
    return x + 1;
}

template <typename FunctionType>
static int doWork()
{
    std::array<FunctionType, 24> functions;

    for (auto& f : functions)
        f = addOne;

    int sum = 0;
    for (auto& f : functions)
        sum += f (4);

    return sum;
}
```

Benchmark against std::function

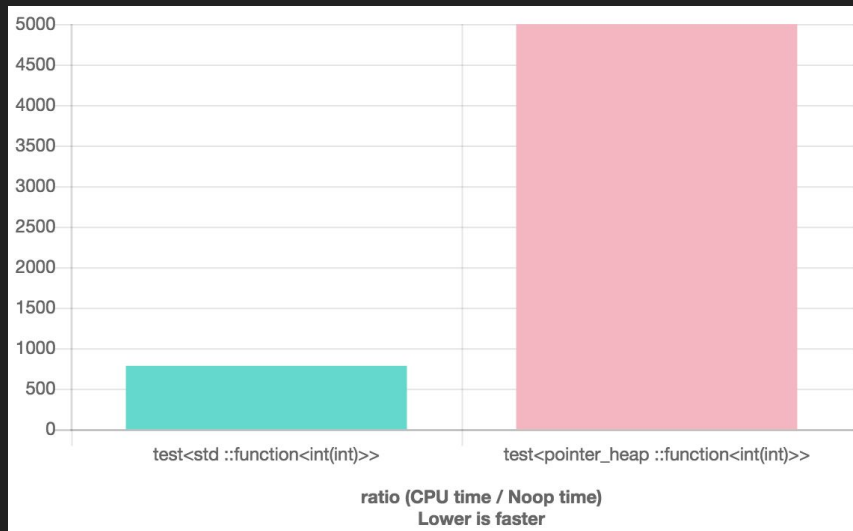
```
int addOne (int x)
{
    return x + 1;
}

template <typename FunctionType>
static int doWork()
{
    std::array<FunctionType, 24> functions;

    for (auto& f : functions)
        f = addOne;

    int sum = 0;
    for (auto& f : functions)
        sum += f (4);

    return sum;
}
```



Courtesy of <http://quick-bench.com/>

You can run my code yourself (see final slide).

The small buffer optimisation (SBO)

The `std::function` in all the common standard libraries contains a small area of memory in which it can store small callable things.

- Good: avoids allocation
- Bad: non-configurable implementation defined size

But we still want this for compatibility in JUCE!

std::function feature parity

```
template <typename Functor>
function (Functor f)
    : invokePtr (reinterpret_cast<invokePtr_t> (invoke<Functor>)),
      createPtr (reinterpret_cast<createPtr_t> (create<Functor>)),
      destroyPtr (reinterpret_cast<destroyPtr_t> (destroy<Functor>))
{
    if (sizeof (Functor) <= sizeof (stack))
    {
        storagePtr = std::addressof (stack);
    }
    else
    {
        heapSize = sizeof (Functor);
        storagePtr = std::malloc (heapSize);
    }

    createPtr (storagePtr, std::addressof (f));
}
```

```
typename std::aligned_storage<24>::type stack;
int heapSize;
void* storagePtr = nullptr;
```

```
function& operator= (function const& other)
{
    if (storagePtr != nullptr)
    {
        destroyPtr (storagePtr);

        if (storagePtr != std::addressof (stack))
            std::free (storagePtr);

        storagePtr = nullptr;
    }

    if (other.storagePtr != nullptr)
    {
        invokePtr = other.invokePtr;
        createPtr = other.createPtr;
        destroyPtr = other.destroyPtr;

        if (other.storagePtr == std::addressof (other.stack))
        {
            storagePtr = std::addressof (stack);
        }
        else
        {
            heapSize = other.heapSize;
            storagePtr = std::malloc (heapSize);
        }

        createPtr (storagePtr, other.storagePtr);
    }

    return *this;
}
```

Benchmark against std::function

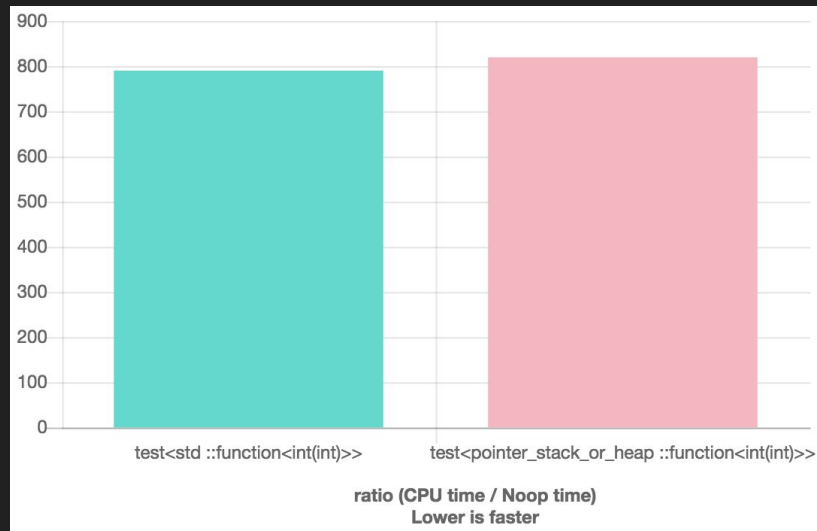
```
int addOne (int x)
{
    return x + 1;
}

template <typename FunctionType>
static int doWork()
{
    std::array<FunctionType, 24> functions;

    for (auto& f : functions)
        f = addOne;

    int sum = 0;
    for (auto& f : functions)
        sum += f (4);

    return sum;
}
```



Courtesy of <http://quick-bench.com/>

Type-erasing function pointer members (again)

```
template <typename Functor>
static Result invoke (Functor* f, Arguments&&... args)
{
    return (*f)(std::forward<Arguments>(args)...);
}

template <typename Functor>
static void create (Functor* destination, Functor* source)
{
    new (destination) Functor (*source);
}

template <typename Functor>
static void destroy (Functor* f)
{
    f->~Functor();
}
```

```
using invokePtr_t = Result (*)(const void*, Arguments&&...);
using createPtr_t = void (*)(void*, const void*);
using destroyPtr_t = void (*)(void*);

invokePtr_t invokePtr;
createPtr_t createPtr;
destroyPtr_t destroyPtr;

typename std::aligned_storage<24>::type stack;
int heapSize;
void* storagePtr = nullptr;
```

Type-erasing with inheritance

```
template <typename ReturnType, typename... Args>
struct FunctorHolderBase
{
    virtual ~FunctorHolderBase() {}
    virtual ReturnType operator()(Args&&...) = 0;
    virtual void copyInto (void*) const = 0;
    virtual FunctorHolderBase<Result, Arguments...>* clone() const = 0;
};
```

```
typename std::aligned_storage<32>::type stack;
FunctorHolderBase<Result, Arguments...>* functorHolderPtr = nullptr;
```

```
template <typename Functor, typename ReturnType, typename... Args>
struct FunctorHolder final : FunctorHolderBase<Result, Arguments...>
{
    FunctorHolder (Functor func) : f (func) {}

    ReturnType operator()(Args&&... args) override
    {
        return f (std::forward<Arguments> (args)...);
    }

    void copyInto (void* destination) const override
    {
        new (destination) FunctorHolder (f);
    }

    FunctorHolderBase<Result, Arguments...>* clone() const override
    {
        return new FunctorHolder (f);
    }

    Functor f;
};
```

Benchmark against std::function

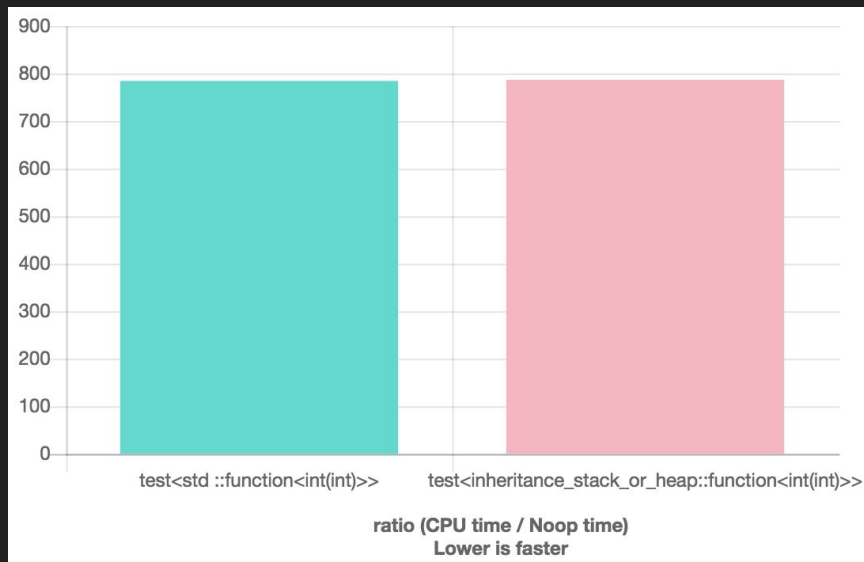
```
int addOne (int x)
{
    return x + 1;
}

template <typename FunctionType>
static int doWork()
{
    std::array<FunctionType, 24> functions;

    for (auto& f : functions)
        f = addOne;

    int sum = 0;
    for (auto& f : functions)
        sum += f (4);

    return sum;
}
```



Courtesy of <http://quick-bench.com/>

Replacing std::function

```
namespace std {  
    template <typename T>  
    using function = juce::function<T>;  
}
```

Adding declarations to the std namespace is undefined behaviour (apart from some template specialisations)

...

but adding `function` (carefully guarded by lots of preprocessor checks to make sure we are not going to collide with a preexisting definition) has not caused any issues that we know about.

- Only done on systems which lack a std::function implementation
- JUCE is distributed as source code
- No ABI issues

Simple stack function

```
template <typename Functor>
function (Functor f)
{
    static_assert (sizeof (FunctorHolder<Functor, Result, Arguments...>) <= sizeof (stack), "Too big!");
    functorHolderPtr = (FunctorHolderBase<Result, Arguments...>*) std::addressof (stack);
    new (functorHolderPtr) FunctorHolder<Functor, Result, Arguments...> (f);
}
```

```
Result operator() (Arguments&&... args) const
{
    return (*functorHolderPtr) (std::forward<Arguments>(args)...);
}
```

```
typename std::aligned_storage<24>::type stack;
FunctorHolderBase<Result, Arguments...>* functorHolderPtr = nullptr;
```

Benchmark against std::function

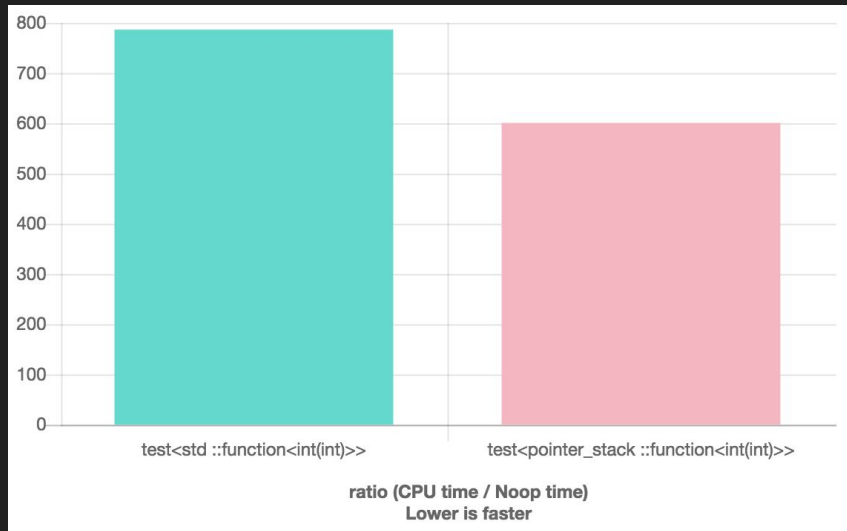
```
int addOne (int x)
{
    return x + 1;
}

template <typename FunctionType>
static int doWork()
{
    std::array<FunctionType, 24> functions;

    for (auto& f : functions)
        f = addOne;

    int sum = 0;
    for (auto& f : functions)
        sum += f (4);

    return sum;
}
```



Courtesy of <http://quick-bench.com/>

Memory management

```
template <typename, class Allocator = std::allocator<char>>
class function;

template <typename Result, typename... Arguments, class Allocator>
class function<Result (Arguments...), Allocator>
{
public:
    template <typename Functor>
    function (Functor f, const Allocator& alloc = Allocator())
```

Delegate allocation behaviour to an allocator

- Allocators are awkward in C++11, better ones available in C++17

A better approach: Louis Dionne's 2017 CppCon talk - "Runtime Polymorphism: Back to the Basics"

More indirection

```
template <typename>
class function;

template <typename Result, typename... Arguments>
class function<Result (Arguments...)>
{
public:
    virtual ~function() {}
    virtual Result operator() (Arguments&&...) const = 0;
};

template <typename, size_t>
class StackFunction;

template <size_t stackSize, typename Result, typename... Arguments>
class StackFunction<Result (Arguments...), stackSize> final : public function<Result (Arguments...)>
```

Non-type erasing function

```
template <typename Functor>
function (Functor f)
    : functionPtr (f)
{}

function() = default;

Result operator() (Arguments&&... args) const
{
    return functionPtr (std::forward<Arguments>(args)...);
}

Result(*functionPtr)(Arguments...) = nullptr;
```

Benchmark against std::function

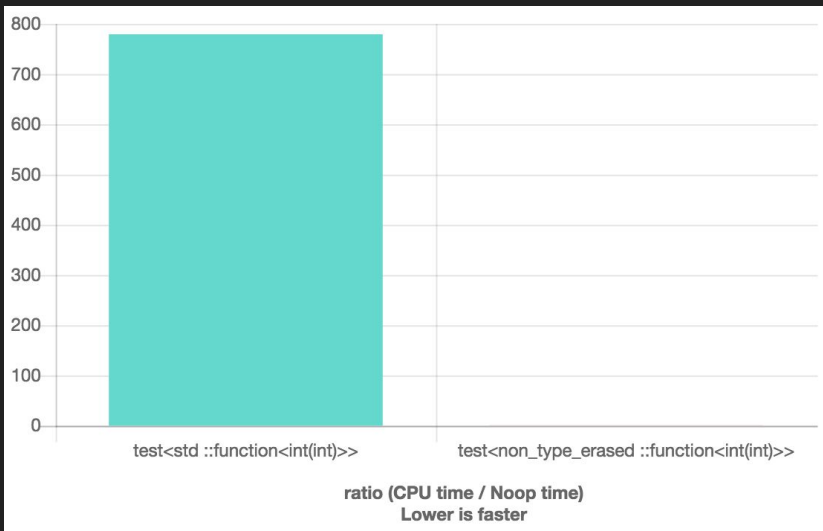
```
int addOne (int x)
{
    return x + 1;
}

template <typename FunctionType>
static int doWork()
{
    std::array<FunctionType, 24> functions;

    for (auto& f : functions)
        f = addOne;

    int sum = 0;
    for (auto& f : functions)
        sum += f (4);

    return sum;
}
```



Courtesy of <http://quick-bench.com/>

Caveats

Only use a type-erased callable wrapper if:

- You don't know the memory layout of the callables you need to store
- You need to store different callables in the same container

Capturing function arguments with a templated type and calling that type is often a better approach.

Conclusions

Rolling your own version of `std::function` is not difficult

Replace the internal memory management with your own scheme

- Enforce realtime safety at compile time
- Increase runtime performance
- Use more adventurous memory management

If you can restrict your callable objects to those with a certain memory layout then do not use a type-erasing wrapper!

Links

Code from the slides: <https://github.com/tpoole/presentations>

JUCE: <https://juce.com/>

SG14 inplace_function: <https://github.com/WG21-SG14/SG14>

folly::Function: <https://github.com/facebook/folly>

folly::Function CppCon talk: <https://www.youtube.com/watch?v=SToaMS3jNH0>

Quick Bench: <http://quick-bench.com/>

Allocation strategies: <https://www.youtube.com/watch?v=gVGtNFg4ay0>